

Botzer_AI879_HW_Q2_Week7

February 5, 2024

```
[ ]: # Author: Brandon Botzer  
# Date: 2/03/2024  
# Class: Penn State - AI 879
```

Q2:

Implement the Adaptive non-maximal suppression (ANMS) (chapter 7.1.1) and test it on an image from the class or an image of your choice. The code should take as input an image and a desired number of points. It should return a list of coordinates for the interest points similar to the Harris algorithm. Below is a pseudo-code for the ANMS algorithm:

- 1 Read image
- 2 Detect corners
- 3 Sort corners by intensity metric in descending order
- 4 Initialize a radius vector (rad) with same length as the number of corners, and set the radius for first element to infinity (or large enough number)
- 5 Loop corners from second to the last one (note current corner as c1)
 - 5.1. Loop corners from first corner to current corner (note current corner as c2)
 - 5.1.1. Calculate euclidean distance $tmp_dist = euclidean(c1, c2)$
 - 5.2. End Loop
 - 5.3 Assign the radius for c1 as $rad(c1) = \min(tmp_dist)$ over inner loop
 - 5.4. End Loop
- 6 Sort “rad” by value descending
- 7 Keep top N corners
- 8 Display image and corners

```
[ ]: # Imports for functions  
  
# The scikit-image package provides a wide variety of filter applications  
# which reduce the need to write out the corr / conv matrices  
  
from PIL import Image
```

```

import skimage as ski
import numpy as np
import matplotlib.pyplot as plt
import cv2 as cv

# Fast sorting and computations
import scipy as sci
from scipy.spatial.distance import cdist
# Showing the results of scipy computations
import pandas as pd

cv.__version__, ski.__version__

```

```
[ ]: ('4.8.1', '0.22.0')
```

```

[ ]: # Read in the image
image = ski.io.imread('L07 greatvalley1.jpg') # I had to find a low-res image
↳for this
# Convert to grayscale
image = ski.color.rgb2gray(image)

# IF NOT USING A LOW-RES IMAGE AND YOU NEED TO REDUCE THE IMAGE SIZE, Uncomment
↳the rescale below

# After some initial trials, my initial image is too high-res for this
# the amount of corners being found is ~230,000 and requires too much memory
↳to process
# I am going to rescale the image to something smaller or find a low-res image

# image = ski.transform.rescale(image, 0.5, order=1)

# Number of points to use for ANMS
n_points = 50

```

1 The function to run the ANMS

```

[ ]: def calcANMS(image, n_pts=50):

    # Change a color image to grayscale if it has not been done
    if len(image.shape) > 2:
        image = ski.color.rgb2gray(image)

    # Compute the Harris_response image to find peaks (corners)
    harris_response = ski.feature.corner_harris(image)

```

```

# Find the corners. Set min_distance=1 to find all possible peaks
# Setting min_distnace = larger number effectivly implements the ANMS but
↪ I'll do it manually here
# and compare results later to corner_peaks().
# ski.feature.corner_peaks(): https://scikit-image.org/docs/stable/api/skimage.feature.html#skimage.feature.corner\_peaks
↪ skimage.feature.html#skimage.feature.corner_peaks
harris_corners = ski.feature.corner_peaks(harris_response, min_distance=1)

# Create the radius vector with same length as number of corners and set
↪ the first element to inf
# np.full_like() will allow an array to be built with the same shape as the
↪ harris_response
rad_vector = np.zeros_like(harris_corners)

# This is the faster scipy method with clever matrix setup

# Calculate the Euclidian distance between each corner point to each other
↪ corner point (creates a square matrix)
dist = cdist(harris_corners, harris_corners)
# Set the same corner to corner distance to infinity along the identity line
# (we don't want zero here since we're not counting the corner against
↪ itself)
np.fill_diagonal(dist, np.inf)

# Show the distances so we know what we're doing
dist_df = pd.DataFrame(dist)
# print('The corner distance matrix:\n', dist_df)

# Take the minimum of each row to find the shortest radius vector for each
↪ corner-to-corner
rad_vec = np.min(dist, axis=1)

# Show the first 10 radius vectors for sanity
# print(rad_vec[0:10])

# Sort radius vectors by distance (descending order)
# There may be a better way to do this but this works well and is fast
↪ enough

# Storage list for the intensities
rad_list = []

# For each of the corners
for r_v in rad_vec:

```

```

    # Append to the list the intensity at the corner
    rad_list.append(r_v)

    ind = np.argsort(rad_list) # Sorts ascending and returns the indecies of
    the sort
    # Flip for decending
    ind = np.flip(ind)

    # Assign the sorted indexies to the harris_corner and rad_vector to sort
    the arrays
    sorted_harris_corners = harris_corners[ind]
    sorted_rad_vec = rad_vec[ind]

    # Create a copy of the origional image that we're going to modify to show
    the ANMS items
    marked_image = np.copy(image)

    # Convert the grayscale back to a 3 channel color so we can plot red circles
    marked_image = ski.color.gray2rgb(marked_image)

    # Take the tuple from the harris corners and pull them apart for a later
    for loop
    sorted_x_centers = sorted_harris_corners[0:,0]
    sorted_y_centers = sorted_harris_corners[0:,1]

    fig, ax = plt.subplots(ncols=1, nrows=1)

    # Count the number of circlces that will occur outside of the image
    skipped = 0

    # for n_points worth of x,y points with a radius, create a circle
    for x_pt, y_pt, rad in zip(sorted_x_centers[:n_pts], sorted_y_centers[:
    n_pts], sorted_rad_vec[:n_pts]):
        # This provides back in row, column
        c_y, c_x = ski.draw.circle_perimeter(int(y_pt), int(x_pt), int(rad))#,
        shape=marked_image.shape)

        # Mark up the image with a red circle
        # Note: the circle is [R,G,B] in float values {0,1}

        # If any location on the circlce radius is outside of the image, do not
        plot that circlce.
        if np.any(c_y>=800) or np.any(c_x>=450):
            skipped += 1
        else:

```

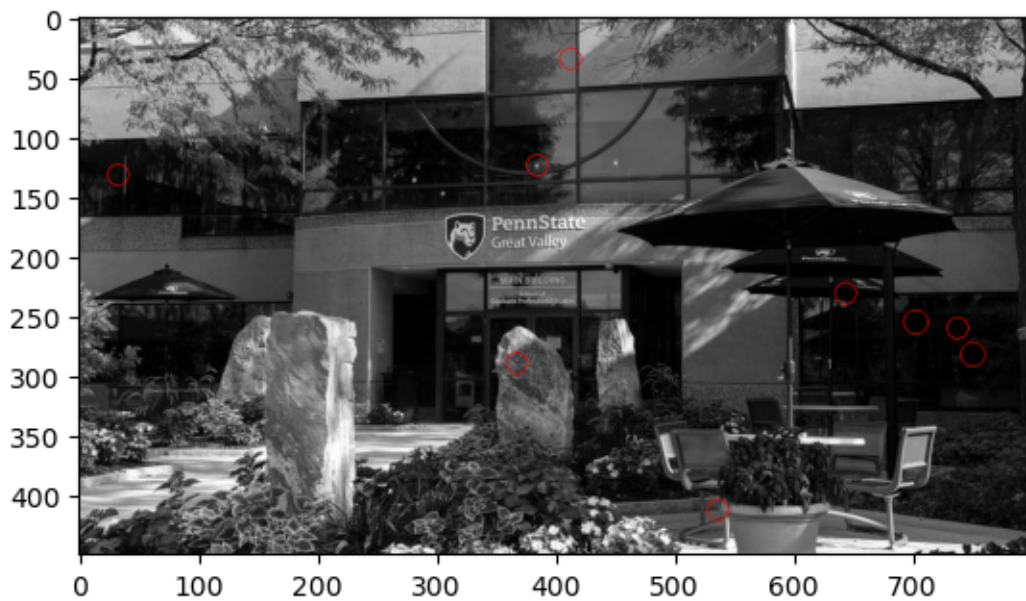
```
marked_image[c_x, c_y] = (1,0,0) # the row, column from
↪circle_perimeter needs flipped again for plotting
```

```
# Display the marked up image
ax.imshow(marked_image)
plt.show()

print(f'Number of points skipped: {skipped}')

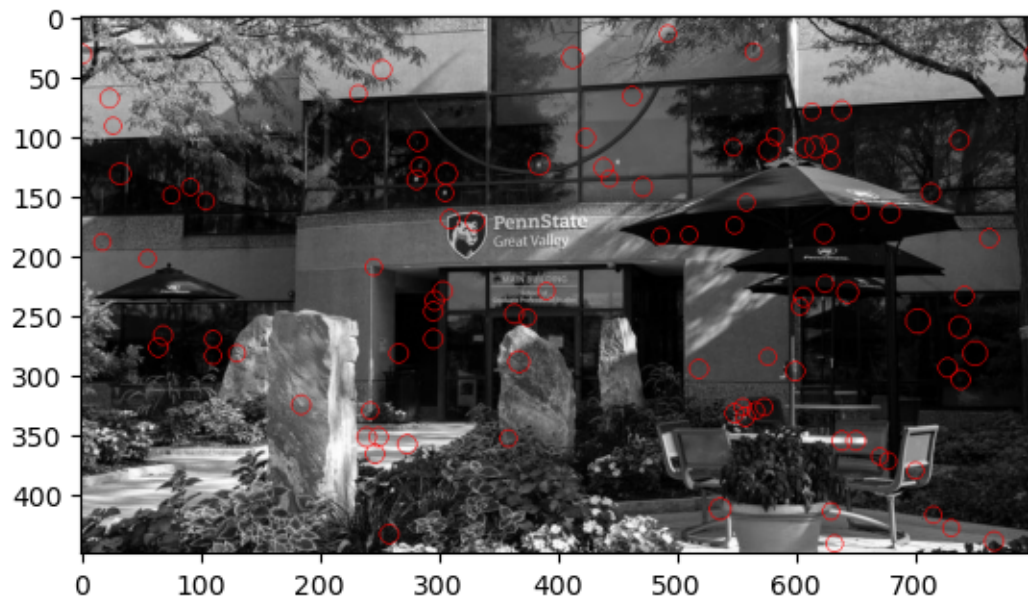
return sorted_harris_corners, sorted_rad_vec
```

```
[ ]: corners, rad_vector = calcANMS(image, n_pts=10)
```



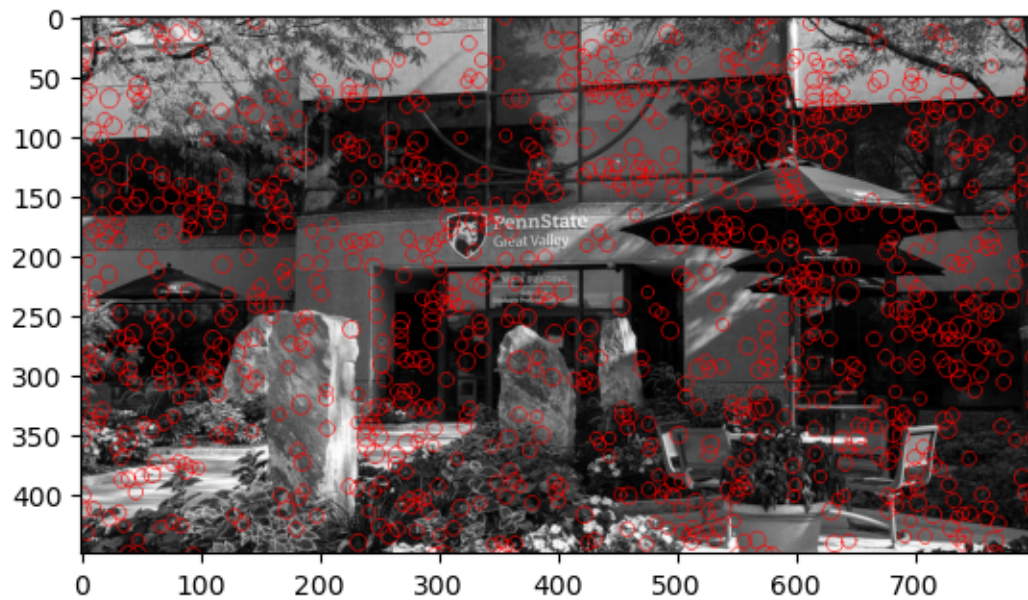
Number of points skipped: 1

```
[ ]: corners, rad_vector = calcANMS(image, n_pts=100)
```



Number of points skipped: 1

```
[ ]: corners, rad_vector = calcANMS(image, n_pts=1000)
```



Number of points skipped: 35

2 A further look at each step in the function

This explains how I arrived at using `cdist` rather than a double for loop with simpler sorting

```
[ ]: # Compute the Harris_response image to find peaks (corners)
harris_response = ski.feature.corner_harris(image)

# Find the corners. Set min_distance=1 to find all possible peaks
# Setting min_distance = larger number effectively implements the ANMS but I'll
↳do it manually here
# and compare results later to corner_peaks().
# ski.feature.corner_peaks(): https://scikit-image.org/docs/stable/api/skimage.
↳feature.html#skimage.feature.corner_peaks
harris_corners = ski.feature.corner_peaks(harris_response, min_distance=1)

[ ]: len(harris_corners)

[ ]: 8811

[ ]: harris_corners[0:10]

[ ]: array([[448, 743],
          [448, 299],
          [448, 750],
          [447, 426],
          [448, 505],
          [448, 754],
          [348, 576],
          [448, 708],
          [443, 287],
          [390, 231]], dtype=int64)
```

This sorting of the Harris corners no longer needs to be done.

By using the `cdist` vectorization on the radii, I am able to find the corner to corner distance for each corner quickly. Then I sort these distances and return the index of the sort (which is inverted for descending).

By applying this index to both the radius vector and the Harris corners, I have sorted both in one move.

```
[ ]: # Sort Harris corners by intensity (descending order)
# There may be a better way to do this but this works well

'''
# Storage list for the intensities
intensity_list = []

# For each of the corners
```



```

for h_corner in harris_corners:
    # Append to the list the intensity at the corner
    intensity_list.append(image[h_corner[0], h_corner[1]])

ind = np.argsort(intensity_list) # Sorts ascending and returns the indecies of
    ↪ the sort
# Flip for decending
ind = np.flip(ind)

# Assign the sorted indexies to the harris_corner to sort the array
sorted_harris_corners = harris_corners[ind]
'''

```

```

[ ]: '\n# Storage list for the intensities\nintensity_list = []\n\n# For each of the
corners\nfor h_corner in harris_corners:\n    # Append to the list the intensity
at the corner\n    intensity_list.append(image[h_corner[0], h_corner[1]])\n\nind
= np.argsort(intensity_list) # Sorts ascending and returns the indecies of the
sort\n# Flip for decending\nind = np.flip(ind)\n\n# Assign the sorted indexies
to the harris_corner to sort the array\nsorted_harris_corners =
harris_corners[ind]\n'

```

```

[ ]: # Create the radius vector with same length as number of corners and set the
    ↪ first element to inf
# np.full_like() will allow an array to be built with the same shape as the
    ↪ harris_response
rad_vector = np.zeros_like(harris_corners)

# Don't have to do this due to the vectorized cdist and collapsing functions
# Assign the first element to a large value
# rad_vector[0][0] = 900000
# rad_vector[0][1] = 900000
# #Check the first element
# print(rad_vector[0], rad_vector[1])

```

This is the method described via the class notes. It is extreemly long in computation time due to the nested loops.

It also requies an initial sorting of the harris corners, and then a secondary sorting through the radius vectors to help with the speed.

It is extreemly slow... I believe for my low-res image it took around 10 minutes.

```

[ ]: # Corner looping - This is incredibly slow as it needs to go
# through nested for loops. There is a faster method...
# using scipy's cdist function and some clever matrix setups
# By applying the cdist, I am able to forgo the multiple sorts needed
# and can apply them both later.

```



```
'''
# From 2nd corner to the last
for i, c1 in enumerate(sorted_harris_corners[1:]):
    tmp_dist = 900000000
# From the first corner to the current corner
    for j, c2 in enumerate(sorted_harris_corners[0:i]):

        # Calculate the Euclidian distance using numpy linear algebra norms
        dist = np.linalg.norm(c2-c1)
        # Take the minimum of the distances
        tmp_dist = min(tmp_dist, dist)

# Assign the minimum radius to the ith rad_vector
rad_vector[i] = tmp_dist
'''
```

```
[ ]: '\n# From 2nd corner to the last\nfor i, c1 in
enumerate(sorted_harris_corners[1:]):\n    tmp_dist = 900000000\n# From the
first corner to the current corner\n    for j, c2 in
enumerate(sorted_harris_corners[0:i]):\n\n        # Calculate the Euclidian
distance using numpy linear algebra norms\n            dist =
np.linalg.norm(c2-c1)\n            # Take the minimum of the distances\n
tmp_dist = min(tmp_dist, dist)\n\n        # Assign the minimum radius to the ith
rad_vector\n    rad_vector[i] = tmp_dist\n'
```

Instead of doing two sorts, I can implement `cdist` from `scipy`. This will allow the radius vector to be calculated for every corner against every corner.

From here, I can `argsort()` the `rad_vector` and directly apply the indices returned to the `harris_corners` and the `rad_vector` to put them in matched descending order.

```
[ ]: # Maybe use a scipy method cdist: https://docs.scipy.org/doc/scipy/reference/
generated/scipy.spatial.distance.cdist.html

# This is the faster scipy method with clever matrix setup

# Calculate the Euclidian distance between each corner point to each other
↳ corner point (creates a square matrix)
dist = cdist(harris_corners, harris_corners)
# Set the same corner to corner distance to infinity along the identity line
# (we don't want zero here since we're not counting the corner against itself)
np.fill_diagonal(dist, np.inf)

# Show the distances so we know what we're doing
dist_df = pd.DataFrame(dist)
dist_df
```

```

[ ]:
      0      1      2      3      4      5      \
0      inf  444.000000  7.000000  317.001577  238.000000  11.000000
1  444.000000      inf  451.000000  127.003937  206.000000  455.000000
2    7.000000  451.000000      inf  324.001543  245.000000  4.000000
3  317.001577  127.003937  324.001543      inf  79.006329  328.001524
4  238.000000  206.000000  245.000000  79.006329      inf  249.000000
...      ...      ...      ...      ...      ...
8806 107.205410  423.467826  109.562767  301.424949  228.973798  111.085553
8807  86.377080  421.849499  89.050547  297.941269  223.215143  90.785461
8808 166.207701  350.201371  171.125685  234.446156  171.840042  174.000000
8809  94.556861  424.646912  96.932967  301.438219  227.554389  98.488578
8810 204.002451  240.002083  211.002370  113.000000  34.014703  215.002326

      6      7      8      9      ...      8801  \
0  194.650970  35.000000  456.027411  515.274684  ...  564.272984
1  294.497878  409.000000  13.000000  89.375612  ...  400.424774
2  200.688814  42.000000  463.026997  522.230792  ...  569.256533
3  179.724790  282.001773  139.057542  203.160035  ...  405.583530
4  122.641755  203.000000  218.057332  280.071419  ...  429.706877
...      ...      ...      ...      ...      ...
8806 134.014925  102.019606  433.979262  481.016632  ...  471.492312
8807 138.311243  81.154174  432.726241  482.548443  ...  487.225820
8808  55.713553  144.222051  359.869421  401.812145  ...  398.131888
8809 138.361844  90.199778  435.377997  484.058881  ...  482.187723
8810 105.688221  169.002959  252.031744  313.229947  ...  442.747106

      8802      8803      8804      8805      8806      8807  \
0  592.426367  567.810708  489.902031  108.632408  107.205410  86.377080
1  439.209517  405.394869  104.403065  416.433668  423.467826  421.849499
2  597.175016  572.763476  496.762519  111.400180  109.562767  89.050547
3  443.829922  410.478989  189.678676  294.497878  301.424949  297.941269
4  466.061155  434.342031  261.090023  222.272355  228.973798  223.215143
...      ...      ...      ...      ...      ...
8806 496.904417  474.647237  447.017897  7.071068      inf  21.213203
8807 513.454964  490.493629  450.320997  22.360680  21.213203      inf
8808 426.279251  401.639142  365.662413  77.369245  83.952367  93.520051
8809 507.986220  485.396745  451.070948  15.556349  12.649111  9.055385
8810 478.029288  447.235956  292.549141  192.083315  198.600101  191.509791

      8808      8809      8810
0  166.207701  94.556861  204.002451
1  350.201371  424.646912  240.002083
2  171.125685  96.932967  211.002370
3  234.446156  301.438219  113.000000
4  171.840042  227.554389  34.014703
...      ...      ...
8806  83.952367  12.649111  198.600101

```

```

8807    93.520051    9.055385   191.509791
8808           inf    91.082380   148.600135
8809    91.082380           inf   196.331353
8810   148.600135   196.331353           inf

```

```
[8811 rows x 8811 columns]
```

```

[ ]: # Take the minimum of each row to find the shortest radius vector for each
      ↪corner-to-corner
rad_vec = np.min(dist, axis=1)

# Show the first 10 radius vectors for sanity
rad_vec[0:10]

```

```

[ ]: array([6.          , 5.65685425, 4.          , 6.          , 3.16227766,
          4.          , 6.32455532, 6.40312424, 4.24264069, 4.47213595])

```

```

[ ]: # Sort radius vectors by distance (descending order)
      # There may be a better way to do this but this works well and is fast enough

# Storage list for the intensities
rad_list = []

# For each of the corners
for r_v in rad_vec:
    # Append to the list the intensity at the corner
    rad_list.append(r_v)

ind = np.argsort(rad_list) # Sorts ascending and returns the indecies of the
      ↪sort
# Flip for decending
ind = np.flip(ind)

# Assign the sorted indexies to the harris_corner and rad_vector to sort the
      ↪arrays
sorted_harris_corners = harris_corners[ind]

sorted_rad_vec = rad_vec[ind]

```

```

[ ]: # Create a copy of the origional image that we're going to modify to show the
      ↪ANMS items
marked_image = np.copy(image)

# Convert the grayscale back to a 3 channel color so we can plot red circles
marked_image = ski.color.gray2rgb(marked_image)

```

```
[ ]: # Take the tuple from the harris corners and pull them apart for a later for
      ↪ loop
sorted_x_centers = sorted_harris_corners[0:,0]
sorted_y_centers = sorted_harris_corners[0:,1]

[ ]: # Make circles at each of the 'n' radii for on the image

fig, ax = plt.subplots(ncols=1, nrows=1)

# Count the number of circlces that will occur outside of the image
skipped = 0

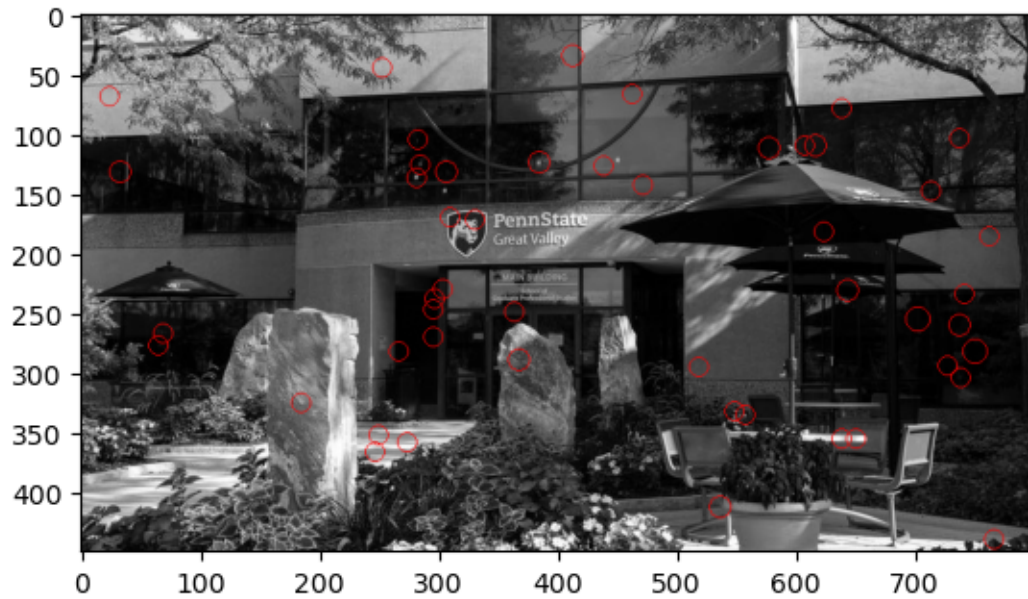
# for n_points worth of x,y points with a radius, create a circle
for x_pt, y_pt, rad in zip(sorted_x_centers[:n_points], sorted_y_centers[:
      ↪ n_points], sorted_rad_vec[:n_points]):
    # This provides back in row, column
    c_y, c_x = ski.draw.circle_perimeter(int(y_pt), int(x_pt), int(rad))#,
      ↪ shape=marked_image.shape)

    # Mark up the image with a red circle
    # Note: the circle is [R,G,B] in float values {0,1}

    # If any location on the circlce radius is outside of the image, do not plot
      ↪ that circlce.
    if np.any(c_y>=800) or np.any(c_x>=450):
        skipped += 1
    else:
        marked_image[c_x, c_y] = (1,0,0) # the row, column from
      ↪ circle_perimeter needs flipped again for plotting

# Display the marked up image
ax.imshow(marked_image)
plt.show()

print(f'Number of points skipped: {skipped}')
```



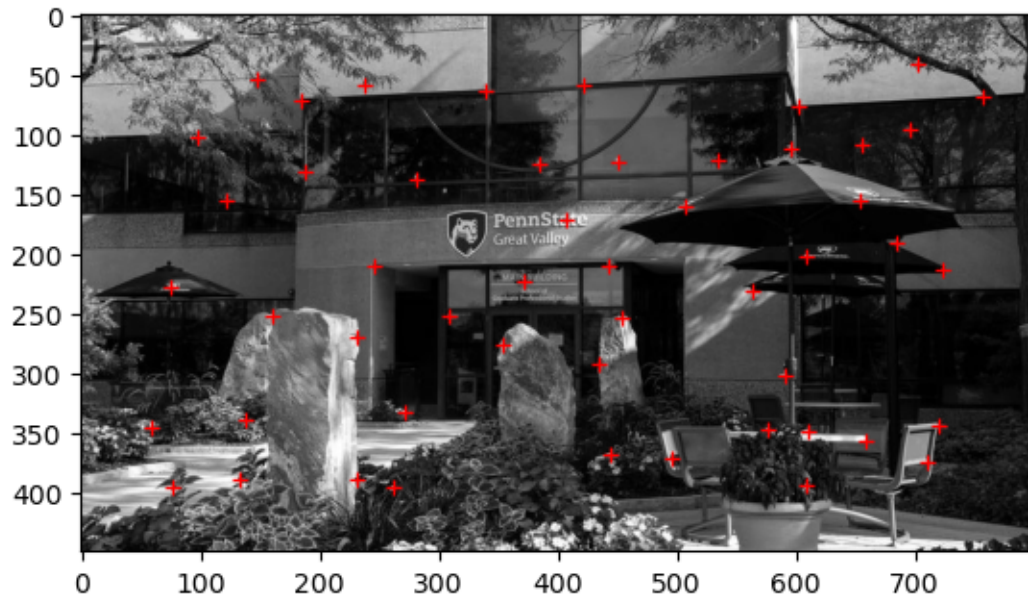
Number of points skipped: 1

I am going to try to replicate my results from ANMS with just the skimage library

```
[ ]: hcs = ski.feature.corner_harris(image)

[ ]: peaks = ski.feature.corner_peaks(hcs, min_distance=30)

[ ]: fig, ax = plt.subplots()
      ax.imshow(image, cmap='gray')
      ax.plot(peaks[:, 1], peaks[:,0], color='red', marker='+', linestyle='None')
      plt.show()
```

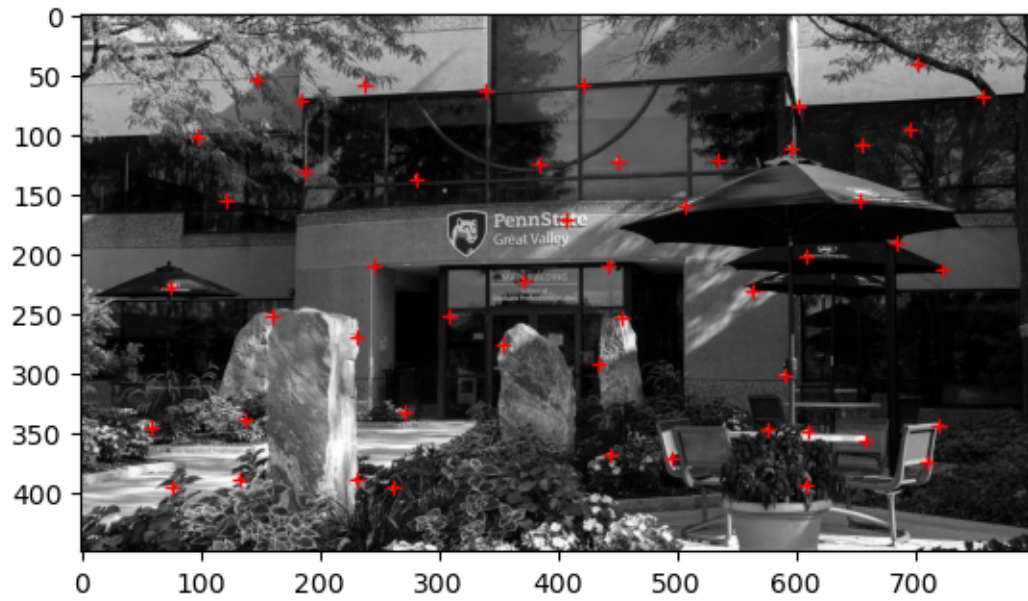


`corner_peaks` differs from `skimage.feature.peak_local_max` in that it suppresses multiple connected peaks with the same accumulator value.

See: https://scikit-image.org/docs/stable/api/skimage.feature.html#skimage.feature.corner_peaks

```
[ ]: local_peaks = ski.feature.peak_local_max(hcs, min_distance=30)

[ ]: fig, ax = plt.subplots()
      ax.imshow(image, cmap='gray')
      ax.plot(local_peaks[:, 1], local_peaks[:, 0], color='red', marker='+',
              ↪linestyle='None')
      plt.show()
```



Notice that the `min_distance` does not correspond exactly to what ANMS is doing. ANMS is finding the peaks with the largest radii and plotting them. `corner_peaks` and `peak_local_max` are finding peaks that are separated by some minimum distance from other peaks at a certain threshold. This would be similar (though not exactly like) plotting ANMS and then removing radii that overlapped.